

Machine Learning Theory, Self-Study

Selected Solutions for Mohri's Foundation to ML (2nd)

Others to be added on an ad hoc basis

To the curious: ML is not the same thing as deep learning. Theory here does not include neural network work and requires more human intervention by nature. The latter will be included in its own repo with PyTorch scripts. Algorithmic analysis and probabilistic considerations found in CLRS text work under dev.

Guilherme Albertini

Contents

1	PAC Learning	1
1.1	Priors	1
1.2	Exercises	1

Preface

TODO blah blah

Chapter 1

PAC Learning

1.1 Priors

The appendix is quite helpful for bounds used throughout the text. As the authors state, “The book of Kearns and Vazirani (1994) is an excellent reference dealing with most aspects of PAC-learning and several other foundational questions in machine learning. Our example of learning axis-aligned rectangles, also discussed in that reference, is originally due to Blumer et al. (1989).”

Mohri’s course notes will be helpful: <https://cs.nyu.edu/~mohri/ml20/>. Note that the textbook pdf is freely available on his website, too. Familiarity with data science fundamentals, multivariable calculus, linear algebra, and algorithmic analysis is assumed for this specific text. Knowledge of convex optimization (or nonlinear programming) and real analysis would be very useful, it seems.

I include a “corrected” and (imo) clearer proof from other authors that do not assume continuity of distribution under the chapters folder: http://compbio.fmph.uniba.sk/vyuka/ml/handouts/rectangles_correction.pdf

For proofs throughout: For implication $p \implies q$, an antecedent (or hypothesis) p is a sufficient condition for a consequent (a conclusion) q when the truth of p alone implies the truth of q ; however, p being false does not always imply that q is also false. A necessary condition is when the truth of q is guaranteed by the truth of p , or we can say that the truth of p is implied by the truth of q ; in other words, p is not possible without q . Several necessary conditions may induce a condition whereas a sufficient condition is alone enough to produce the said condition. The sufficient term is the part that immediately follows “if” and the necessary term is the part that immediately follows the “then”. Note that these are converses, but the converse may not always be true. Further reading: <https://philosophy.stackexchange.com/questions/22/what-is-the-difference-between-necessary-and-sufficient>, <https://www.kaptest.com/study/lSAT/lSAT-formal-logic-necessary-vs-sufficient/>

1.2 Exercises

Exercise 2.1

Necessary and sufficient: a concept class \mathcal{C} is efficiently PAC-learnable using hypothesis space \mathcal{H} in the standard PAC model if and only if it is efficiently PAC-learnable using the hypothesis space $\mathcal{H} \cup \{h_0, h_1\}$ in the two-oracle PAC model.

Sufficiency: Show that if \mathcal{H} is PAC-learnable in the standard, one-oracle model then so too is it in the variant.

Since \mathcal{C} is efficiently PAC-learnable using \mathcal{H} , there exists an algorithm \mathcal{A} and a polynomial p such that for any distribution \mathcal{D} and any target concept $c \in \mathcal{C}$, if \mathcal{A} is given a sample of size $m \geq p(1/\epsilon, 1/\delta, n, \text{size}(c))$ drawn from \mathcal{D} , it outputs a hypothesis h from \mathcal{H} such that with probability at least $1 - \delta$, the error of h on \mathcal{D} is at most ϵ . Assume a distribution \mathcal{D} on $\mathcal{X} \times \{-1, +1\}$ where both oracles generate samples from

$\mathcal{D}_+ \propto \mathcal{D}(x, 1)$ and $\mathcal{D}_- \propto \mathcal{D}(x, -1)$. The learner's goal is to output a hypothesis with such probability over the choice of two training sets (in the stochastic scenario where the output label is a probabilistic function of the input, thus does not guarantee unique labels) and requires both $\mathbb{P}[R(h)_{x \sim \mathcal{D}_+} \leq \epsilon] \geq 1 - \delta$ and $\mathbb{P}[R(h)_{x \sim \mathcal{D}_-} \leq \epsilon] \geq 1 - \delta$. By the law of total probability for the error rate (or risk),

$$\begin{aligned} R(h)_{\mathcal{D}} &= \mathbb{E}_{x \sim \mathcal{D}} [\mathbb{1}_{h(x) \neq c(x)}] = \sum_{x \in \mathcal{D}} x \mathbb{P}[\mathbb{1}_{h(x) \neq c(x)}] = \mathbb{P}_{x \sim \mathcal{D}}[h(x) \neq c(x)] \\ &= \mathcal{D}^+(R(h)_{\mathcal{D}_+}) + \mathcal{D}^-(R(h)_{\mathcal{D}_-}) \\ &= \mathcal{D}^+(R(h)_{\mathcal{D}_+}) + (1 - \mathcal{D}^+)(R(h)_{\mathcal{D}_-}) \end{aligned}$$

Now for a weighted sampling method from the negative and positive instance distributions we can say $\epsilon_{\mathcal{D}} = \mathcal{D}^+(\alpha\epsilon)_{\mathcal{D}_+} + \mathcal{D}^-(\beta\epsilon)_{\mathcal{D}_-} = \underbrace{\epsilon(\mathcal{D}^+(\alpha - \beta) + \beta)}_{\leq 1}$ for some constants $0 \leq \alpha, \beta \leq 1$. Note that $\alpha = \beta =$

$0 \implies \epsilon = 0$, which means you're demanding that the learned hypothesis perfectly match the true target function, which can lead to overfitting and complex hypotheses that are computationally expensive to work with (recall $\epsilon > 0$ by definition). Conversely, by setting $\epsilon = 1$, you're ok accepting any hypothesis without regard to its performance. Thus, select an appropriate δ and, for simplicity, set $\mathcal{D}^+ = \frac{1}{2}$, to get

$$\begin{aligned} \mathbb{P}[R(h)_{\mathcal{D}} \leq \frac{\alpha\epsilon}{2}] &\geq 1 - \delta \\ \frac{1}{2}(\mathbb{P}[R(h)_{\mathcal{D}_+} \leq \alpha\epsilon] + \mathbb{P}[R(h)_{\mathcal{D}_-} \leq \alpha\epsilon]) &\geq 1 - \delta \end{aligned}$$

If we set $\alpha = 1$ for this case we then see that by selecting $\mathbb{P}[R(h)_{\mathcal{D}} \leq \frac{\epsilon}{2}]$ with an appropriate confidence interval, we must have that both $\mathbb{P}[R(h)_{\mathcal{D}_-} \leq \epsilon], \mathbb{P}[R(h)_{\mathcal{D}_+} \leq \epsilon] \geq 1 - \delta$. We immediately notice that a biased dataset would then require us to make considerable adjustments to the overall error rate, as \mathcal{D}^+ will shift the weight of distribution of samples. Another thing to note is that by setting α or β to 0, we're essentially requiring that the hypothesis have zero error on positive (or negative) instances (perhaps requiring an absurdly complex hypothesis to do so). This transforms the problem into a rather stringent one-oracle PAC model focused solely on the positive (or negative) class which is not sensitive to noise. The PAC framework is designed to find a balance between minimizing errors on both classes while accounting for uncertainties in real-world data; this decoupled mode tells the learner to query the oracle for instances of one class while imposing stringent error requirements on the other class.

Necessary: Show that if \mathcal{H} is PAC-learnable in the two-oracle variant, it is also PAC-learnable in the standard model.

We now can assume that \mathcal{C} is efficiently PAC-learnable in the two-oracle PAC model so there exists an algorithm \mathcal{A} such that for $c \in \mathcal{C}$, $\epsilon, \delta > 0$, there are m^+ and m^- in $p(1/\epsilon, 1/\delta, \text{size}(c))$, such that if we draw at least this number of negative and positive instances with confidence $1 - \delta$, the hypothesis h output by the learner satisfies:

$$\begin{aligned} R(h)_{\mathcal{D}^{+,-}} &\leq \epsilon \\ \mathbb{P}[R(h)_{\mathcal{D}^{+,-}}] &\leq \mathbb{P}[\epsilon] = \epsilon \end{aligned}$$

So, given sufficient numbers of negative and positive examples, we can generate a hypothesis h such that it has low errors on both negative and positive instances. If we draw too few examples, the conclusions about the hypothesis's performance might not hold true for the entire distribution (generalize); to bridge the gap between the variant and standard model it is then best to take $m \geq \max\{m^+, m^-\}$ and, with conditions above,

$$\mathbb{P}[R(h)_{\mathcal{D}}]$$